

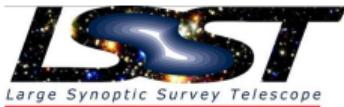
Big data in astronomical observations

William O'Mullane



Data management
LSST
Tucson, AZ USA

3rd July 2019
ISSI
Bern, Switzerland

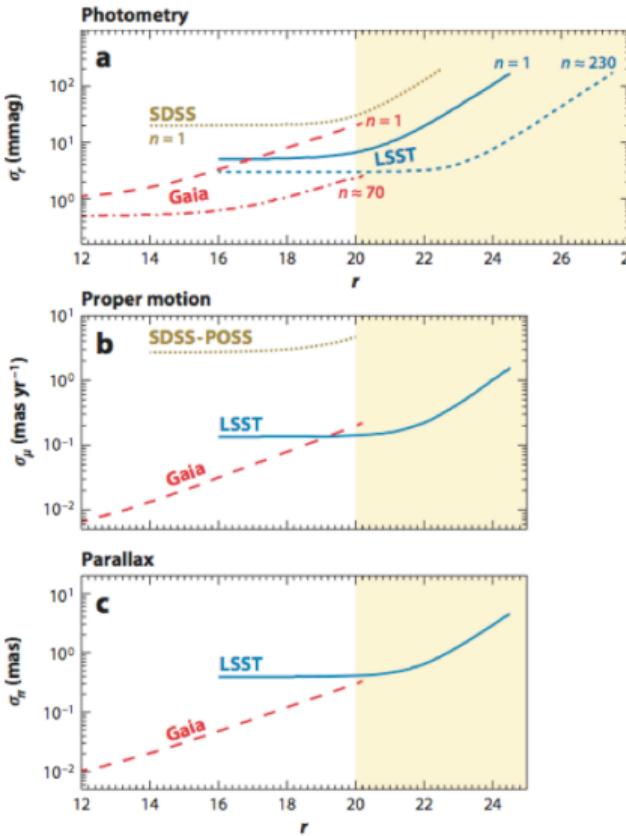




- 1985ish started with BASIC on a Commodore
- 1993 MSc BSc Computer Science, University College Cork, Ireland
- 1993 - 1997 Spacecraft Control Systems (C++), ESA ESOC Germany
- 1997 - 2001 Hipparcos, Integral, Planck, Gaia, Bepi-Sax (C,Java,Oracle, HTM, HEALPix), ESA ESTEC Netherlands
- 2001-2003 Commercial programming - some Astronomy (Java)
- 2003-2005 The Johns Hopkins, SDSS and National Virtual Observatory (C,C#,Java,Sqlserver)
- 2005-2014 Gaia Astrometric Solution and Science Operations (Java, Oracle, Intersystems Cache)
- 2012 PhD in Physics on Implementing the Gaia Astrometric Solution, Barcelona University
- 2014-2017 ESA Division head - all Science Ground Segments in Development
- 2017- LSST Data Management Project Manager (Python,C++), Deputy Project Manager for Software (control systems)

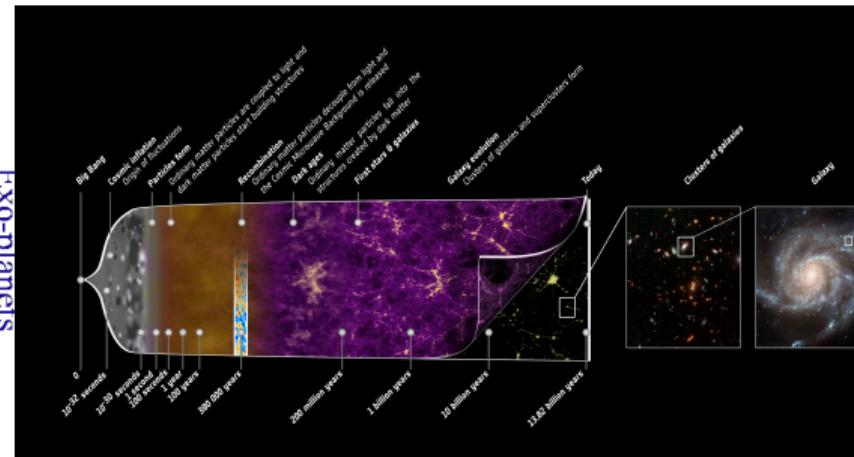
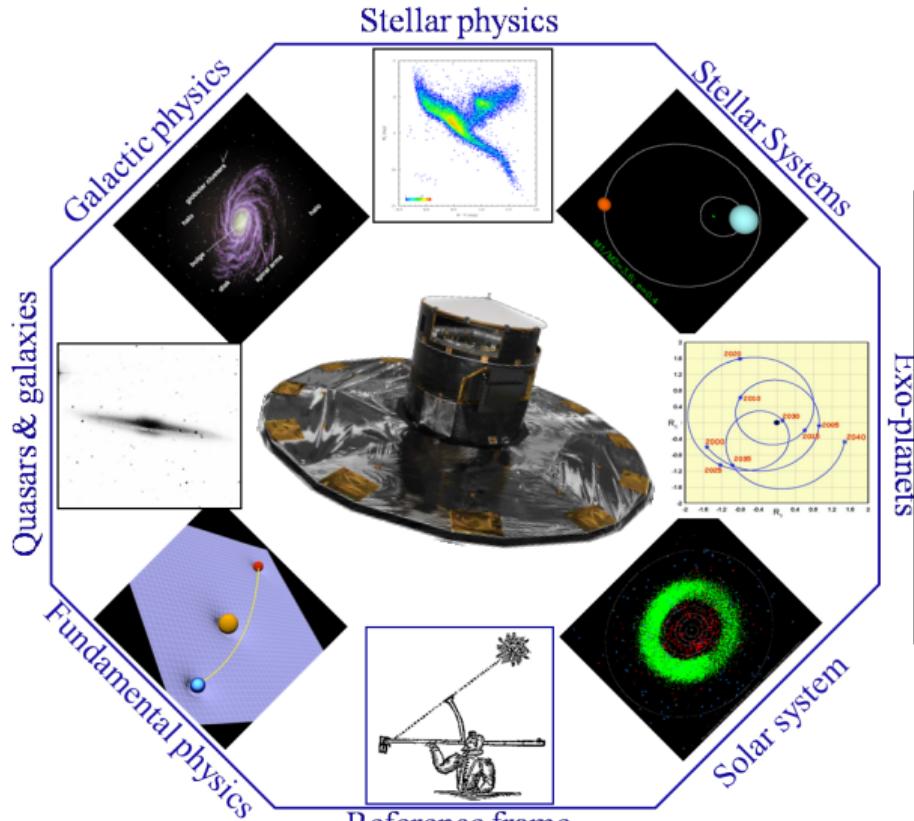


Gaia and LSST continuum



- Gaia: excellent astrometry (and photometry), but only to $r > 20$
- LSST: photometry to $r < 27.5$ and time resolved measurements to $r < 24.5$
- Complementary: photometric, proper motion and trigonometric parallax errors are similar around $r=20$

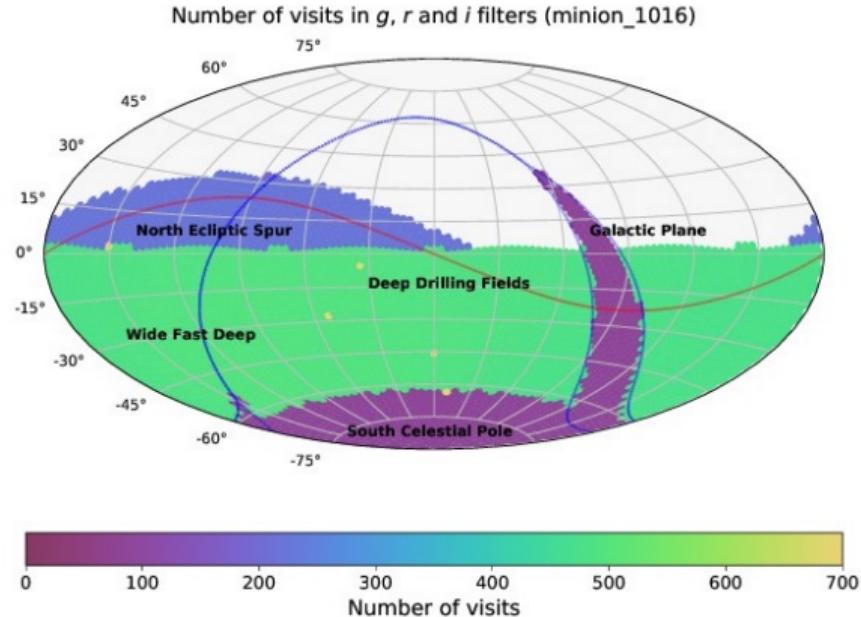
The Milky Way disk belongs to Gaia, and the halo to LSST (plus very faint and/or very red sources, such as white dwarfs and LT(Y) dwarfs). Zeljko Ivezić





An optical/near-IR survey of half the sky in ugrizy bands to r 27.5 (36 nJy) based on 825 visits over a 10-year period: deep wide fast.

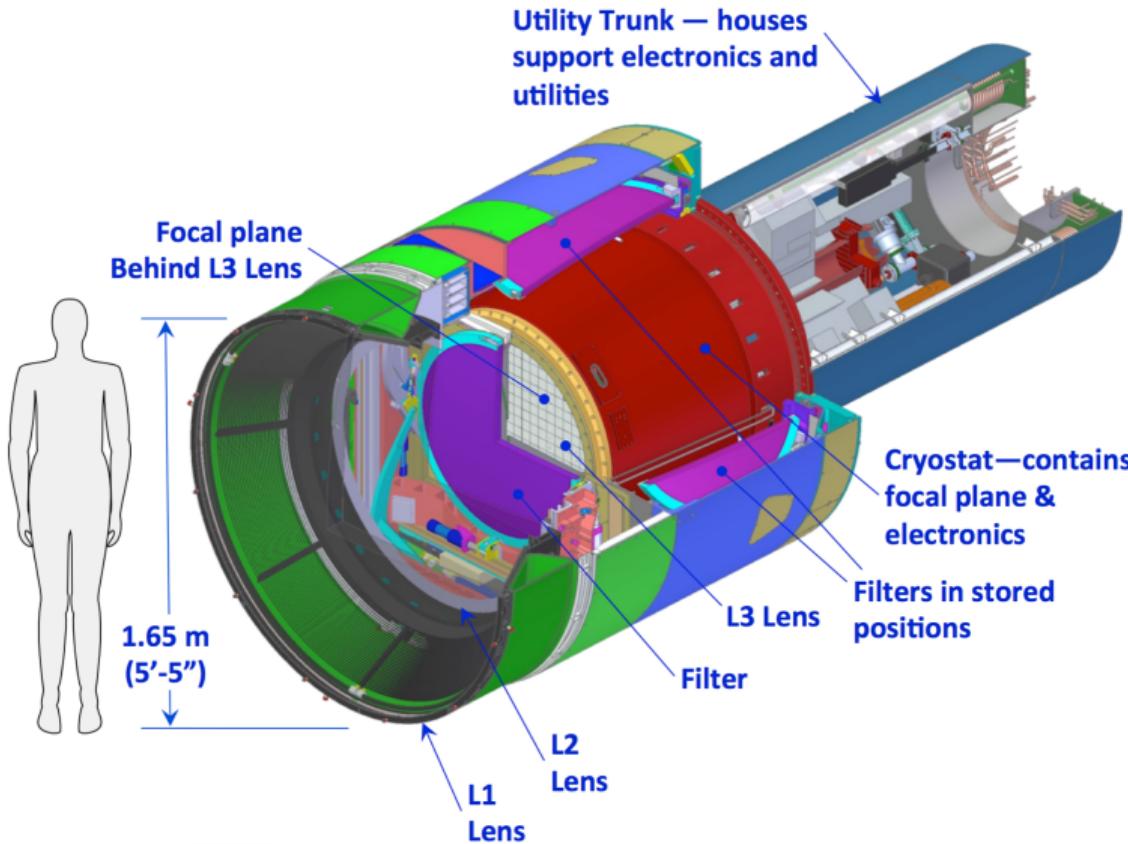
- 90% of time spent on uniform survey: every 3-4 nights, the whole observable sky scanned twice per night
- 100 PB of data: about a billion 16 Mpix images, enabling measurements for **40 billion objects!**



10-year simulation of LSST survey: number of visits in u,g,r band (Aitoff projection of eq. coordinates)

see also <http://www.lsst.org> and Ivezić et al.

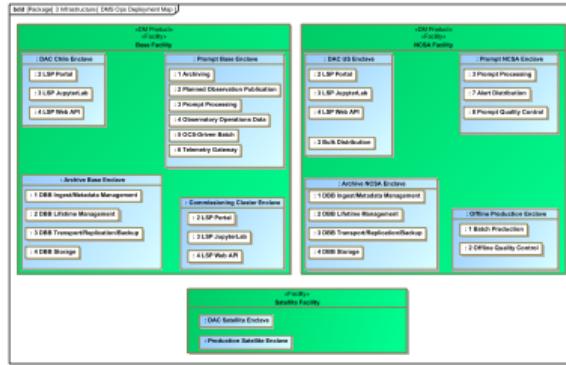
(2008)-arXiv:0805.2366



The largest astronomical camera:

- 2800 kg
- 3.2 Gpix



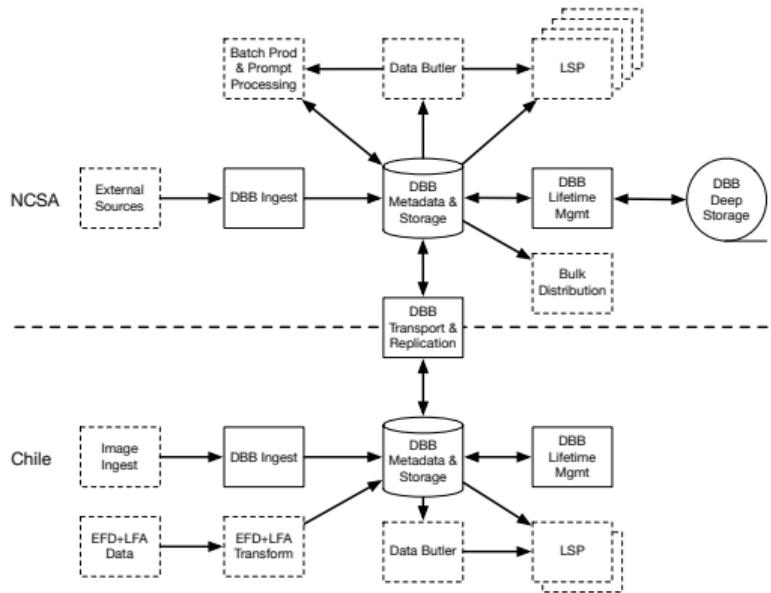


DM must build everything to get LSST products (see <http://ls.st/dpdd>) to the users.

- large data sets (20TB/night)
 - complex analysis
 - aiming for small systematics
 - Science Alerts in under 2 minutes .. (aiming for 1 minute)

About $\frac{1}{2}$ million lines of code (C++/python) all open source on github

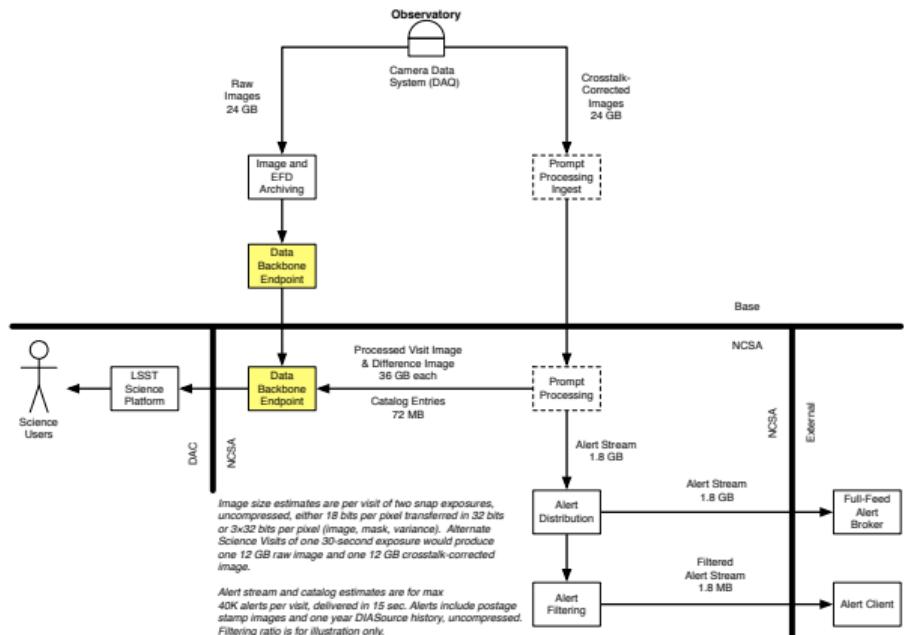
diagram K.T. Lim



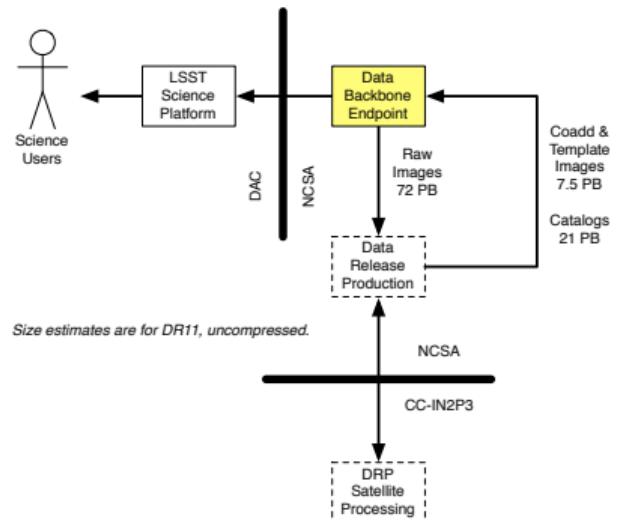
One small box on the previous slide was Data Backbone.

That hides several things

- Qserv - the LSST end user database
 - Custom Massive Parallel Processing (MPP)
 - allow queries on ≈ 20 Petabytes of tabular data
 - 4×10^{10} objects, 4×10^{13} sources (observations)
- All the networks : we now have fiber to the mountain and from La Serena to NCSA (two routes)



Lots to do every night ..
Plus annually there is a data release

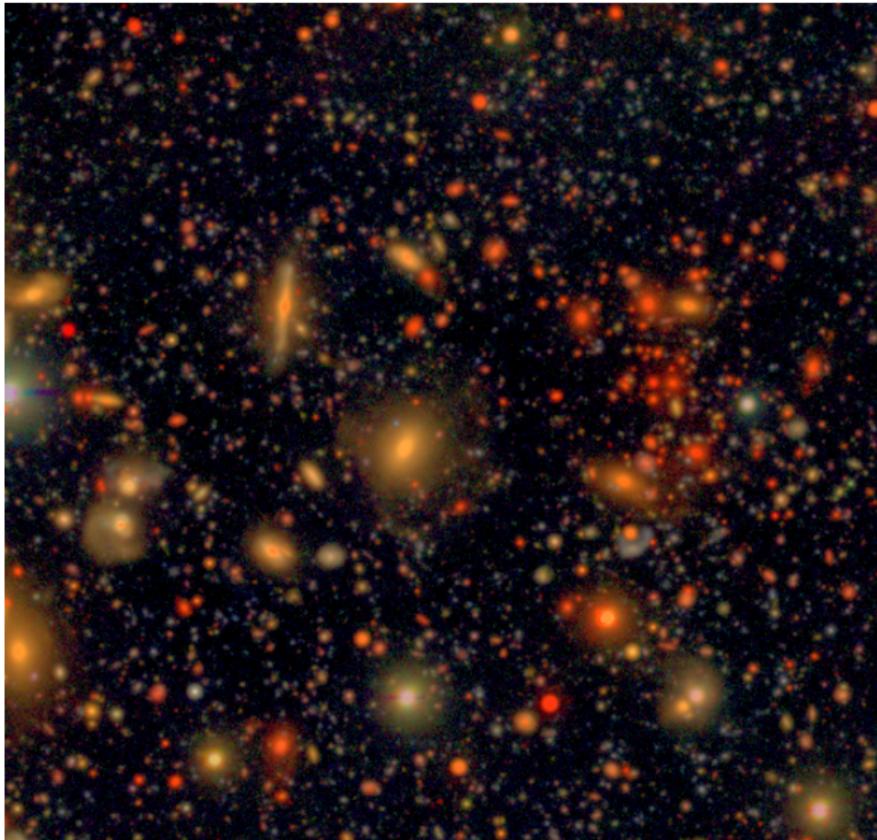


Images from K.T. Lim



Nice colors Lupton et al. (2004)
 $\approx 3.5'$

Image Robert Lupton



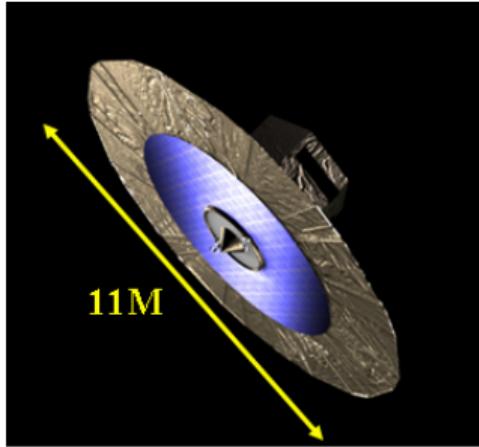
HSC image (COSMOS) from g,r(1.5 hrs) ,i(3 hrs) PSF matched co-add (≈ 27.5)

Challenges:

- Unknown statistical distributions,
Truncated, censored and missing data,
Unreliable quantities (e.g. unknown
systematics and random errors)
- PSF - short exposure - atmosphere
dominated ?
 - cosmic shear signal from weak lensing
- Photometry challenging - will Gaia help
LSST ..
- Everything is blended!!

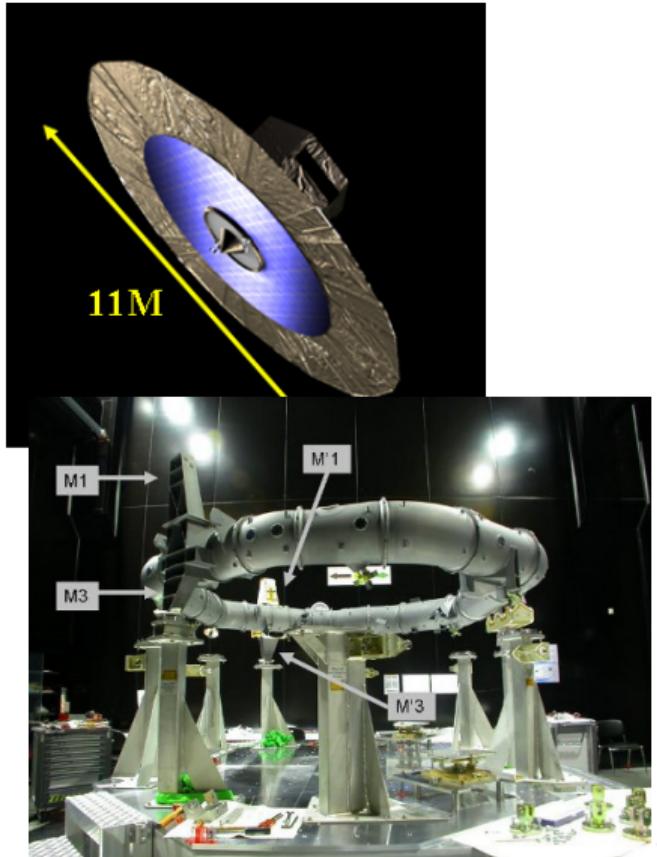


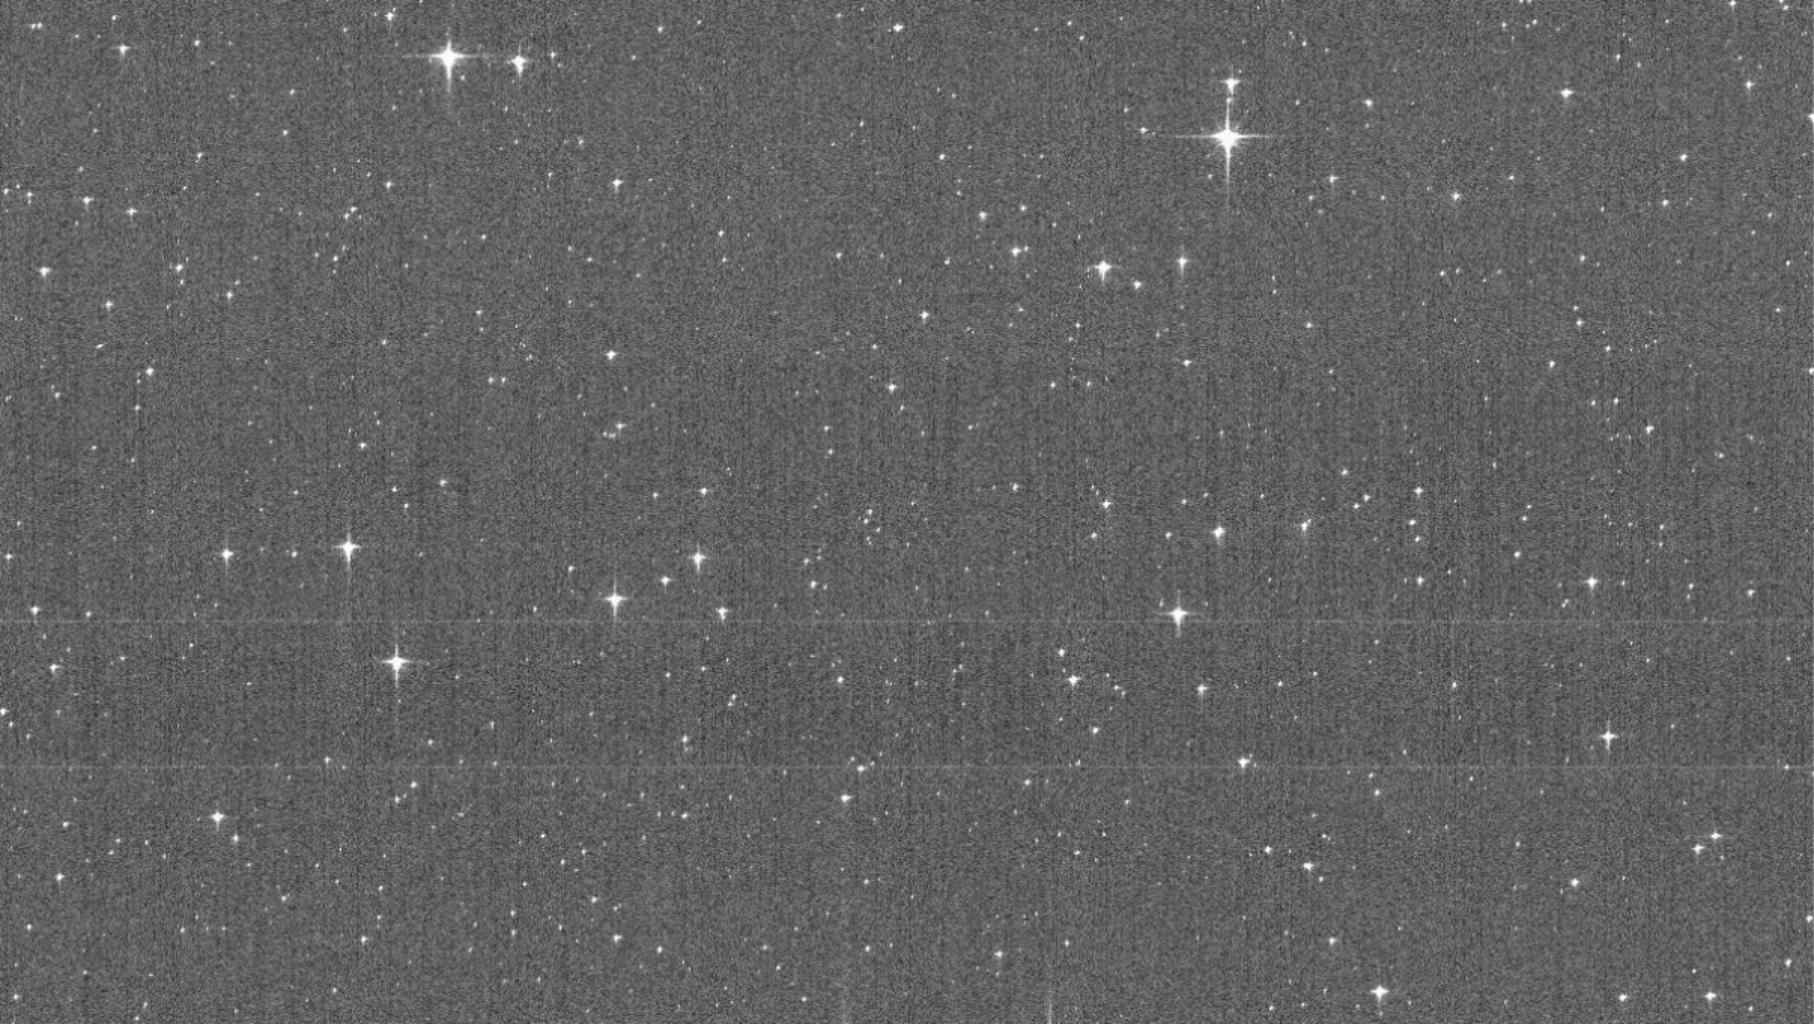
- Mission:
 - ESA Corner Stone 6
 - ESA provided the hardware and launch
 - Mass: 2120 kg (payload 743 kg)
 - Power: 1631 W (payload 815 W)
 - Launched December 19th 2013
 - Stereoscopic Census of Galaxy over 5 years
 - Extended 2 yrs - request for five
 - μ arcsec Astrometry G < 20 (10^9 sources)
 - Radial Velocities G < 16
 - Photometry millimag G < 20

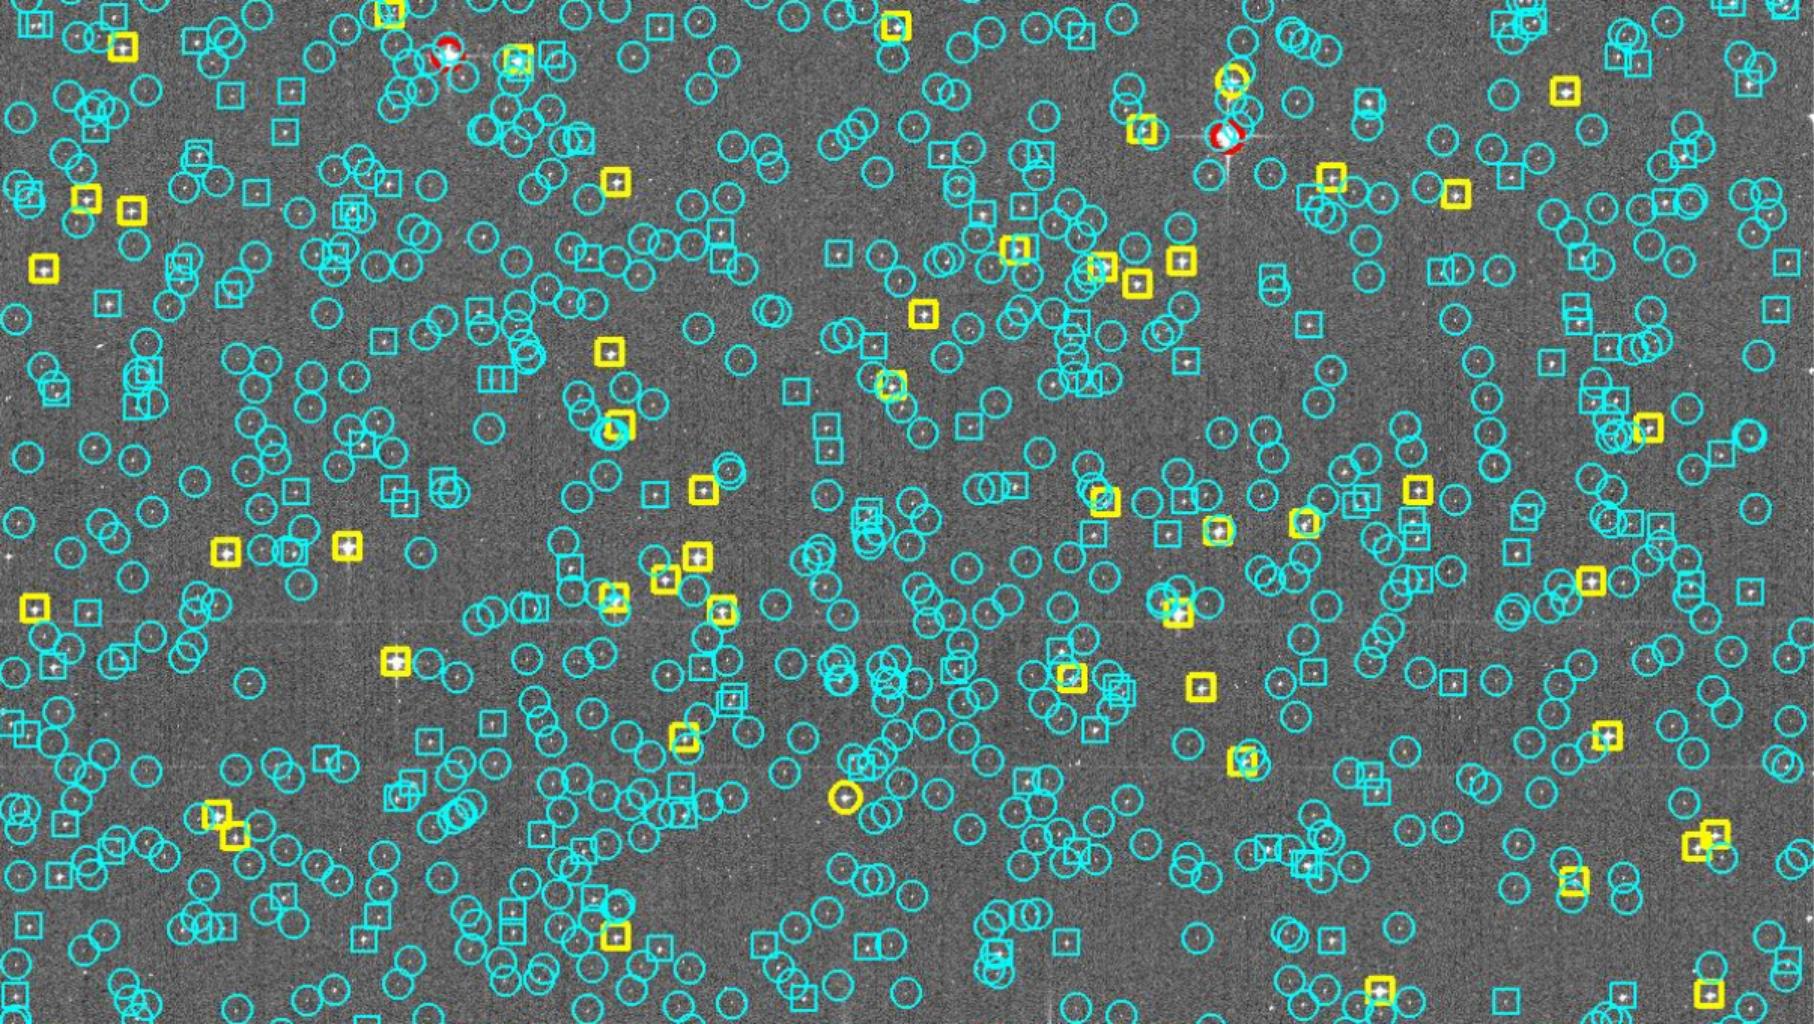




- Mission:
 - ESA Corner Stone 6
 - ESA provided the hardware and launch
 - Mass: 2120 kg (payload 743 kg)
 - Power: 1631 W (payload 815 W)
 - Launched December 19th 2013
 - Stereoscopic Census of Galaxy over 5 years
 - Extended 2 yrs - request for five
 - μ arcsec Astrometry G < 20 (10^9 sources)
 - Radial Velocities G < 16
 - Photometry millimag G < 20
- Final catalogue \approx 202?









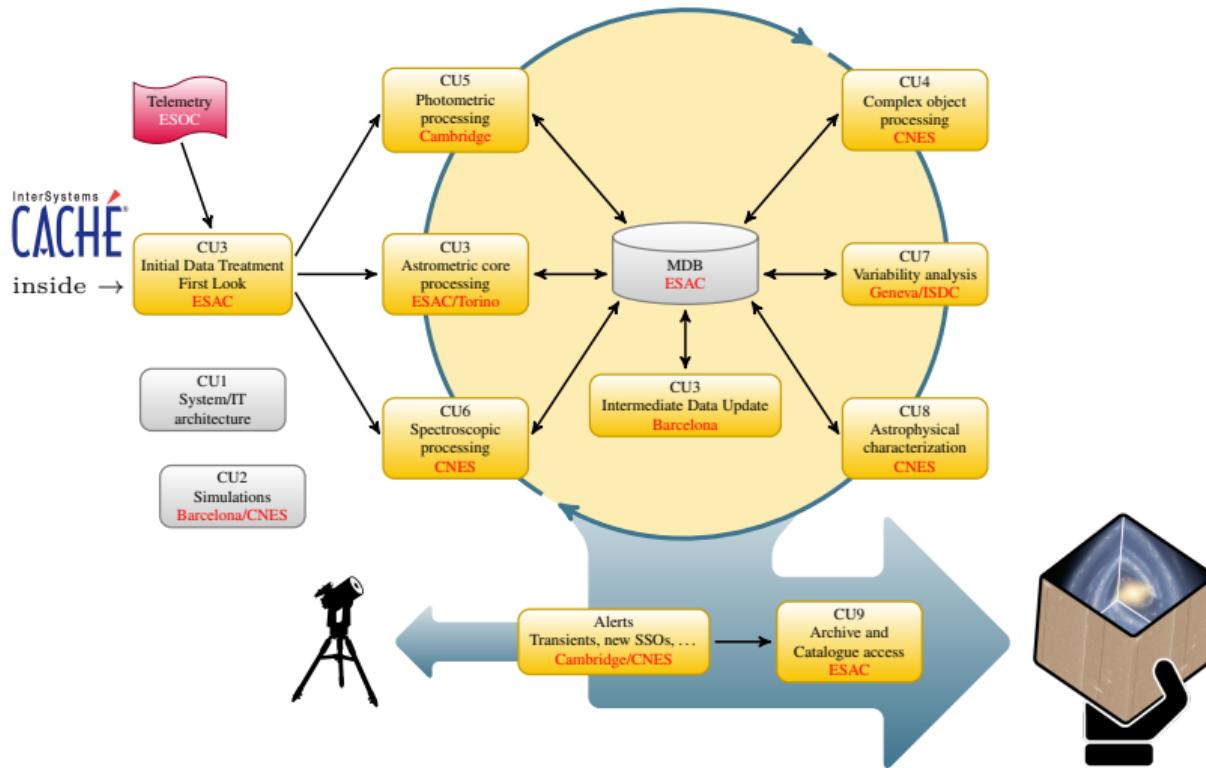
- Astrometric centroid of the CCD image to be determined to an accuracy of 1% of the pixel size!
 - There will be 10^{12} images $\approx 100\text{TB}$ downlink need to handle $\approx 1\text{PB}$
 - At 1 millisecond each that is ≈ 30 years
- Reconstructed attitude is required to order $10 \mu\text{arcsec}$
 - Path of light through instrument needed to nanometre level
 - System must be extremely stable
 - Must consider relativistic light bending from solar system objects.
- Attitude and Geometric calibration can only be done using Gaia's own observational data. (AGIS) (O'Mullane et al., 2011; Lindegren et al., 2012)
 - this requires a significant portion of the data to be processed iteratively



- Astrometric centroid of the CCD image to be determined to an accuracy of 1% of the pixel size!
 - There will be 10^{12} images $\approx 100\text{TB}$ downlink need to handle $\approx 1\text{PB}$
 - At 1 millisecond each that is ≈ 30 years
- Reconstructed attitude is required to order $10 \mu\text{arcsec}$
 - Path of light through instrument needed to nanometre level
 - System must be extremely stable
 - Must consider relativistic light bending from solar system objects.
- Attitude and Geometric calibration can only be done using Gaia's own observational data. (AGIS) (O'Mullane et al., 2011; Lindegren et al., 2012)
 - this requires a significant portion of the data to be processed iteratively



Upstream → Downstream



CU=Coordination Unit

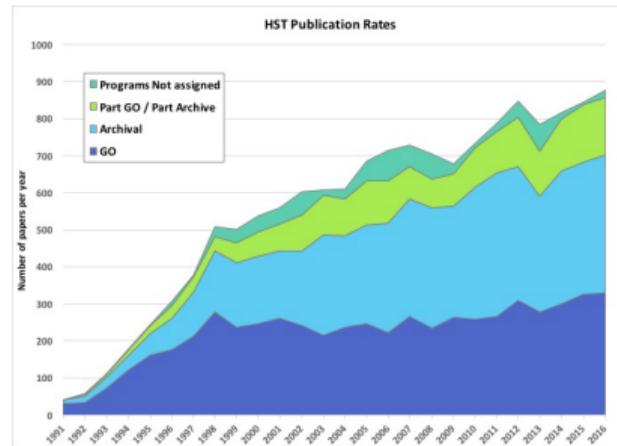
Daily 50 – 100GB
compressed

5yrs 1PB

(see also O'Mullane et al. (2006),
O'Mullane et al. (2009))



- Gaia had 8Mbps downlink - Euclid will have 55Mbps!
- A little like LSST in terms of goals
- Complex ground segment (and consortium)
- Estimated 30PB of data over 6 years
- Launch date 2022



[https://archive.stsci.edu/hst/bibliography/
pubstat.html](https://archive.stsci.edu/hst/bibliography/pubstat.html)

...indicates archival research
probably play an important role in
the scientific success of
XMM-Newton Ness et al. (2014)
Large Synoptic Survey Telescope

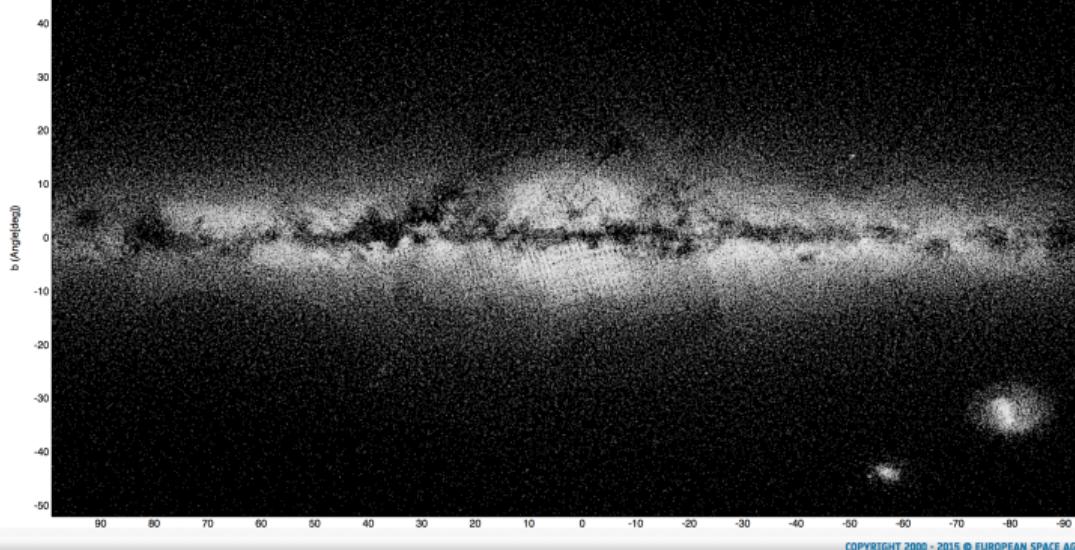


Gaia Archive



EUROPEAN SPACE AGENCY ABOUT ESAC

gaia archive visualization



COPYRIGHT 2000 - 2015 © EUROPEAN SPACE AGENCY

William O'Mullane (womullan)

esa

gaia archive

HOME SEARCH STATISTICS VISUALIZATION HELP DOCUMENTATION VOSPACE SHARE

Simple Form ADQL Form Query Results

Job name: Query examples

Galaxy Data Release 1

- galadr1.allwise_best_neighbour
- galadr1.allwise_neighbourhood
- galadr1.allwise_original_valid
- galadr1.agn_icrf2_match
- galadr1.cephid
- galadr1.ext_phot_zero_point
- galadr1.gala_source
- galadr1.gcc22_best_neighbour
- galadr1.gcc22_neighbourhood
- galadr1.gcc22_original_valid
- galadr1.phot_variable_time_s
- galadr1.phot_variable_time_z
- galadr1.ppmrd_best_neighbour
- galadr1.ppmrd_neighbourhood
- galadr1.ppmrd_original_valid
- galadr1.rnyne
- galadr1.sdss_dr9_best_neighbour
- galadr1.sdss_dr9_neighbourhood

Ctrl+Space for query subcomposition

Reset Form Submit Query

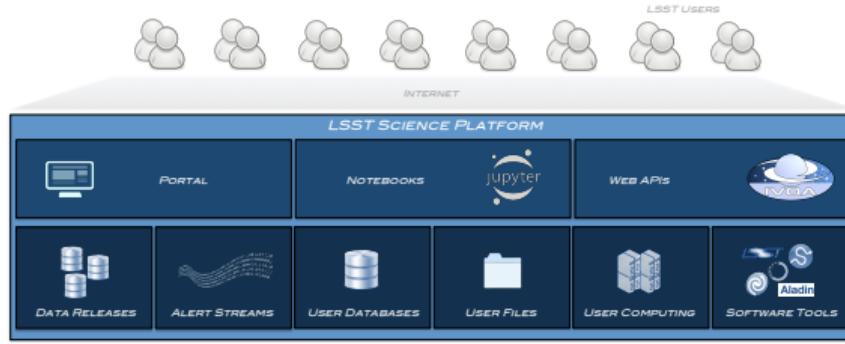
Status	Job	Creation date	Num. rows	Size
<input checked="" type="checkbox"/>	1497719385390	17-Jun-2017, 19:09:45	0 KB	
<input checked="" type="checkbox"/>	outerParallax	05-Nov-2016, 17:27:34	313	10 KB
<input checked="" type="checkbox"/>	14783608081700	05-Nov-2016, 16:46:48	0	0 KB
<input checked="" type="checkbox"/>	14783607846200	05-Nov-2016, 16:46:24	0 KB	
<input checked="" type="checkbox"/>	parallax_diff	04-Nov-2016, 14:47:33	93635	2 MB
<input checked="" type="checkbox"/>	14776481654710	28-Oct-2016, 11:49:25	16285	4 MB
<input checked="" type="checkbox"/>	xmatch_tycho2_agn	10-Jun-2016, 18:10:54	0 KB	

1-13 of 13 Download format: VOTable Apply jobs filter Select all jobs Delete selected jobs

COPYRIGHT 2017 © EUROPEAN SPACE AGENCY. ALL RIGHTS RESERVED. [x1.3]

All Gaia data is publicly accessible at <https://gea.esac.esa.int/archive/>





For DR2:

- Computing: 2,400 cores (≈ 18 TFLOPs)
- File storage: $\approx 4\text{PB}$ (VOSpace)
- Database storage: $\approx 3\text{PB}$ (MYDB)

The Science Platform has three user facing aspects: the Portal (novice), the JupyterLab (intermediate), and the Web APIs (expert and remote tools).

Vision: LSE-319 — Design: LDM-542 — Test: DMTR-51

This is the sort of environment users now expect to have - it is no longer novel. We are finally **Bringing code to the data** - almost did with GAVIP Vagg et al. (2016)

Bruno will have more to say on this topic.



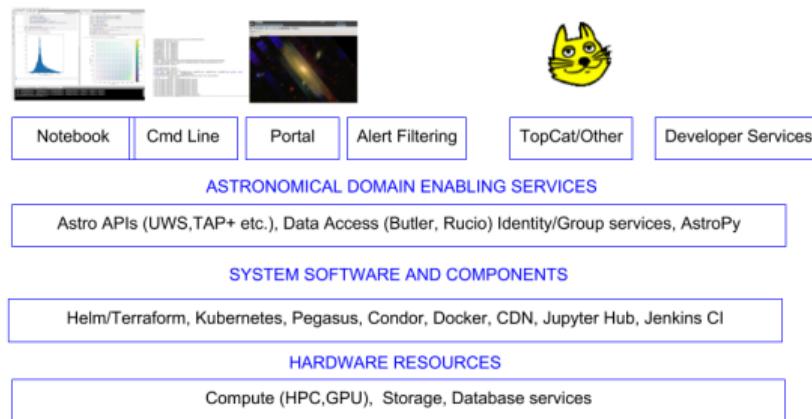
See also Bauer et al. (2019)

Soon, if not already, Data will be looking for astronomers not vice versa.

- Proprietary data may have had its day .. if we want people to look at data we need to remove barriers.
- Networks and infrastructure are improving but more needs to be done
- Who looks after all the data ?
 - For space science in Europe ESAC preserves data - it is relatively small
 - IPAC, HEASARC and STScI pretty much deal with NASA data - still not one location.
 - Who looks after all the ground based data ?
 - There is no long term preservation plan for LSST or other big telescopes in the USA
(Alex wil probably have more to say on this)



We all constantly redesign and rebuild wheels, this will become too expensive as data volume grows ..



- We should agree a component based cyber infrastructure model and work on improving specific components to plug in
 - right now we are all building TAP, Designing Databases , deploying Jupyter ...
- Data models like CAOM from IVOA are going to be essential going forward so we can inter-operate on data

Filesystems are End Of Life - object stores should not be confused with the object databases many of us struggled with in the 90s and 00s .. Google and Amazon do not run filesystems they run object stores.



- LSST allows 9 months for a Data Release Processing Cycle - probably about 6 months actual processing.
- The original requirement for Gaia astrometric solution was three months.
- Traditional batch systems and shared nothing architecture may not always work
 - Gaia Astrometric Solution required temporal spatial access and had global matrices
 - There are processes for LSST which will have similar problems e.g. Forward Global Calibration Model (FGCM) <https://github.com/lsst/fgcmcal> (Burke et al., 2018)
- preparing and staging for tasks can take ever longer as our processing becomes more sophisticated



- In ESA (personal opinion) Gaia was a game changer - the Software was seen as critical
 - I had to fight to keep it off the Launch Critical Items List. We still had a meager 10% or so of the budget.
- LSST recognized early on that the processing was as important as the telescope - DM is one of 4 Subsystem with about 20% of the project budget
- We need to consider long term software support and the software eco system
 - Open source and publicly scrutinized algorithms
 - Agree community cyber infrastructure model (previous slide)
 - Educate astronomers/managers on how to open source
 - Though I made my career in astronomy/computing .. its not easy .. and I really had to do the management thing
- Machine learning:
 - Automated discovery - need it great !
 - Reproducibility, understand ability
- education - everyone needs to be a data scientist/programmer ..



The End



Acronym	Description
AGIS	Astrometric Global Iterative Solution
Archive	The repository for documents required by the NSF to be kept. These include documents related to design and development, construction, integration, test, and operations of the LSST observatory system. The archive is maintained using the enterprise content management system DocuShare, which is accessible through a link on the project website www.project.lsst.org .
CAOM	Common Astronomical Observation Model
CCD	Charge-Coupled Device
CI	Continuous Integration
CU	Coordination Unit
Camera	The LSST subsystem responsible for the 3.2-gigapixel LSST camera, which will take more than 800 panoramic images of the sky every night. SLAC leads a consortium of Department of Energy laboratories to design and build the camera sensors, optics, electronics, cryostat, filters and filter exchange mechanism, and camera control system.
DM	Data Management
DMTR	DM Test Report
Data Backbone	The software that provides for data registration, retrieval, storage, transport, replication, and provenance capabilities that are compatible with the Data Butler. It allows data products to move between Facilities, Enclaves, and DACs by managing caches of files at each endpoint, including persistence to long-term archival storage (e.g. tape).
Data Management	The LSST Subsystem responsible for the Data Management System (DMS), which will capture, store, catalog, and serve the LSST dataset to the scientific community and public. The DM team is responsible for the DMS architecture, applications, middleware, infrastructure, algorithms, and Observatory Network Design. DM is a distributed team working at LSST and partner institutions, with the DM Subsystem Manager located at LSST headquarters in Tucson.



Data Release	The approximately annual reprocessing of all LSST data, and the installation of the resulting data products in the LSST Data Access Centers, which marks the start of the two-year proprietary period.
ESA	European Space Agency
ESAC	European Space Astronomy Centre
ESOC	European Space Operations Centre
ESTEC	European Space Technology Engineering Centre
FGCM	Forward Global Calibration Model
GAVIP	Gaia Added Value Interface Platform
GB	Gigabyte
HEALPix	Hierarchical Equal-Area iso-Latitude Pixelisation
HEASARC	NASA's Archive of Data on Energetic Phenomena
HSC	Hyper Suprime-Cam
HTM	Hierarchical Triangular Mesh
IPAC	No longer an acronym; science and data center at Caltech
IR	Infra Red
ISSI	International Space Science Institute
IVOA	International Virtual-Observatory Alliance
LDM	LSST Data Management (Document Handle)
LSE	LSST Systems Engineering (Document Handle)
LSST	Large Synoptic Survey Telescope
MPP	Massively Parallel Process
NASA	National Aeronautics and Space Administration
NCSA	National Center for Supercomputing Applications
Operations	The 10-year period following construction and commissioning during which the LSST Observatory conducts its survey
PB	PetaByte



PSF	Point Spread Function
Project Manager	The person responsible for exercising leadership and oversight over the entire LSST project; he or she controls schedule, budget, and all contingency funds
Qserv	Proprietary Database built by SLAC for LSST
SDSS	Sloan Digital Sky Survey
Science Platform	A set of integrated web applications and services deployed at the LSST Data Access Centers (DACs) through which the scientific community will access, visualize, and perform next-to-the-data analysis of the LSST data products.
Subsystem	A set of elements comprising a system within the larger LSST system that is responsible for a key technical deliverable of the project.
TAP	Table Access Protocol
TB	TeraByte
XMM	X-ray Multi-mirror Mission (ESA; officially known as XMM-Newton)
arcmin	arcminute minute of arc (unit of angle)
arcsec	arcsecond second of arc (unit of angle)
astrometry	In astronomy, the sub-discipline of astrometry concerns precision measurement of positions (at a reference epoch), and real and apparent motions of astrophysical objects. Real motion means 3-D motions of the object with respect to an inertial reference frame; apparent motions are an artifact of the motion of the Earth. Astrometry per se is sometimes confused with the act of determining a World Coordinate System (WCS), which is a functional characterization of the mapping from pixels in an image or spectrum to world coordinate such as (RA, Dec) or wavelength.
calibration	The process of translating signals produced by a measuring instrument such as a telescope and camera into physical units such as flux, which are used for scientific analysis. Calibration removes most of the contributions to the signal from environmental and instrumental factors, such that only the astronomical component remains.
camera	An imaging device mounted at a telescope focal plane, composed of optics, a shutter, a set of filters, and one or more sensors arranged in a focal plane array.



metric

A measurable quantity which may be tracked. A metric has a name, description, unit, references, and tags (which are used for grouping). A metric is a scalar by definition. See also: aggregate metric, model metric, point metric.



- Bauer, A.E., Bellm, E.C., Bolton, A.S., et al., 2019, arXiv e-prints, arXiv:1905.05116 (arXiv:1905.05116), ADS Link
- [DMTR-51], Bosch, J., Chiang, H.F., Gower, M., et al., 2017, LDM-503-02 (HSC Reprocessing) Test Report, DMTR-51, URL <https://ls.st/DMTR-51>
- Burke, D.L., Rykoff, E.S., Allam, S., et al., 2018, AJ, 155, 41 (arXiv:1706.01542), doi:10.3847/1538-3881/aa9f22, ADS Link
- [LDM-542], Dubois-Felsmann, G., Lim, K.T., Wu, X., et al., 2017, LSST Science Platform Design, LDM-542, URL <https://ls.st/LDM-542>
- Ivezic, Z., et al., 2008, ArXiv e-prints (arXiv:0805.2366), ADS Link
- [LSE-319], Jurić, M., Ciardi, D., Dubois-Felsmann, G., 2017, LSST Science Platform Vision Document, LSE-319, URL <https://ls.st/LSE-319>
- Lindegren, L., Lammers, U., Hobbs, D., et al., 2012, A&A, 538, A78 (arXiv:1112.4139), doi:10.1051/0004-6361/201117905, ADS Link
- Lupton, R., Blanton, M.R., Fekete, G., et al., 2004, PASP, 116, 133 (arXiv:astro-ph/0312483), doi:10.1086/382245, ADS Link
- Ness, J.U., Parmar, A.N., Valencic, L.A., et al., 2014, Astronomische Nachrichten, 335, 210 (arXiv:1311.5751), doi:10.1002/asna.201312001, ADS Link
- O'Mullane, W., Lammers, U., Bailer-Jones, C., et al., 2006, ArXiv Astrophysics e-prints (arXiv:astro-ph/0611885), ADS Link
- O'Mullane, W., Hernández, J., Hoar, J., Lammers, U., 2009, In: D. A. Bohlender, D. Durand, & P. Dowler (ed.) Astronomical Data Analysis Software and Systems XVIII, vol. 411 of Astronomical Society of the Pacific Conference Series, 470, ADS Link
- O'Mullane, W., Lammers, U., Lindegren, L., Hernandez, J., Hobbs, D., 2011, Experimental Astronomy, 31, 215 (arXiv:1108.2206), doi:10.1007/s10686-011-9248-z, ADS Link
- Vagg, D., O'Callaghan, D., O'Hógáin, F., et al., 2016, In: Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series, vol. 9913 of SPIE, 99131V (arXiv:1605.09287), doi:10.1117/12.2233619, ADS Link